







Naturally, we can see we have two Blockwarps in x & y dimensions.

Since each Block is 256 threads wide.
 And
 each warp tile is $\frac{64 \text{ thread wide}}{\text{H.W limitation}}$

$$\begin{aligned} \text{Vector} &= \frac{\text{No. of elements in Warptile}}{64} \\ &= \frac{512}{64} = 8 \end{aligned}$$

Q What is the order in which you decide configuration parameters?

Kernel

Problem

- Data Types of tensors
- Shape
 - ↳ Block Tile
 - Warp Tile
 - Vector
 - Blockwarp

Policy

- How is data going to be mapped b/w threads & global memory.
(Tile distribution)